

Optimization in ML

Machine Learning and Data Mining

Maxim Borisyak

National Research University Higher School of Economics (HSE)

October 2, 2018

Machine Learning in a nutshell

—

ML algorithms

Every (supervised) ML algorithm ever:

- a model \mathcal{A} - set of all possible solutions:

$$\mathcal{A} \subseteq \{f : \mathcal{X} \rightarrow \mathcal{Y}\}$$

- \mathcal{X}, \mathcal{Y} - sample and target spaces.
- a search procedure:

$$S : (\mathcal{X} \times \mathcal{Y})^n \rightarrow \mathcal{A}$$

ML algorithms

Decision Trees:

- model: piece-wise constant functions;
- search procedure: very sinful one.

SVM:

- linear functions (in some space);
- search procedure: margin maximization.

Logistic regression:

- linear functions;
- search procedure: cross-entropy \rightarrow min, any optimization method.

In Deep Learning models are often decoupled from search procedures.

Often, model-search can be factorized further:

- parametrized model:

$$\mathcal{A} = \{f_{\theta} : \mathcal{X} \rightarrow \mathcal{Y} \mid \theta \in \Theta \subseteq \mathbb{R}^n\}$$

- optimization problem:

$$L(f_{\theta}, D) = \mathcal{L}(\theta) \rightarrow_{\theta} \min$$

- $D \in (\mathcal{X} \times \mathcal{Y})^N$ - training set;
- optimization method O :

$$\theta^{t+1} = O[\theta^t, \mathcal{L}].$$

ML algorithms

- parametrized model:
 - heavily domain/data dependent;
- optimization problem:
 - more or less universal:

$$\mathcal{L} = \{\text{log, hinge, } \dots\}\text{-loss} + \{l_1, l_2\}\text{-regularization};$$

- optimization method O :
 - heavily-dependent on nature of \mathcal{L} .

Optimization in ML

Optimization methods:

- heavily restrict:
 - speed of the algorithm;
 - quality of solutions;
- some optimization methods allows for new models:
 - discrete or mixed parameters;
 - variable size models.

Gradient methods

The zoo

- SGD, SGD with momentum:
 - you have no memory;
 - you have to write optimizer in 1 minute;
- Nesterov momentum:
 - you want to fine-tune your solution.
- RMSprop:
 - you have little memory and you gradients explode/vanish;
 - you have 2 minutes before submitting your code for training;
- adagrad, adadelta, adam, adamax:
 - methods to go.

Details are likely to be considered in Deep Learning course.

Second-order methods

Do you have a nearly-quadratic target function?

- yes: is the problem low-dimensional?
 - yes: go Newton!
 - no: use gradient or quasi-Newton methods;
- no: use gradient or quasi-Newton methods.

Hyper-parameter optimization

Hyper-parameter optimization

Hyper-parameter optimization is a meta-algorithm that operates on union of models parametrized by ψ :

$$\mathcal{A} = \bigcup_{\psi} \mathcal{A}_{\psi} = \{f_{\theta_{\psi}}^{\psi} \mid \theta_{\psi} \in \Theta_{\psi}\}$$

- outer loss **might differ** from inner loss:

$$\psi^* = \arg \max_{\psi} Q \left(\arg \min_{\theta_{\psi}} L(\theta_{\psi}) \right)$$

- no sacred meaning, just for convinience:
- example: L - cross-entropy, Q - ROC AUC.

Hyper-parameter optimization

Outer optimization is usually evaluated on a separate set:

- via train-validation-test split;
- via cross-validation;
- the main reason for split into outer and inner problems.

Alternately, BIC or similar can be used.

Outer optimization problem is usually non-differentiable:

- number of units, maximal depth of trees.

Grid-search

Usually, ML algorithms are designed to have **a few, non-sensitive** hyper-parameters:

- outer problem is mostly convex and changes slowly;
- grid-search often works perfectly fine.

Modifications, alternatives:

- randomized grid-search;
- random walk.

Gradient-free methods

Gradient-free methods

- local optimization:
 - 'traditional' methods;
- global optimization:
 - gradient and Hessian are fundamentally local properties;
 - evolutionary methods;
- black-box optimization:
 - variational optimization;
 - Bayesian optimization.

Traditional gradient-free methods

- evaluation of objective function is cheap;
- in practice, **often** can be replaced by gradient-methods:


$$\text{cheap} \Rightarrow \left[\begin{array}{c} \text{closed-form expression} \\ \text{or} \\ \text{allow approximation} \end{array} \right] \Rightarrow \text{differentiable.}$$

- example: Powell, Nelder-Mead.

Multi-start

Just launch local procedure multiple times with different initial guesses.

- each local optima acts like an attractor for local methods;
- effective if depth of local optima positively depend on area of attraction.



Evolution methods

There are just so many...

Basic operations:

- mutation: $x' = x + \text{noise}$;
- crossover: $x' = \text{crossover}(x_1, x_2)$;

Application:

- you have no idea how to optimize objective function;
- evolution algorithms basically can handle any parametrization:
 - e.g. DNA molecules.

Memetic algorithms

```
def memetic(global_step=evolutionary,
            locally_optimize=annealing):

    population = []
    mature_population = [ <random> ]

    while ...:
        population = global_step(mature_population)
        mature_population = [
            locally_optimize(x) for x in population
        ]
    return mature_population
```

```
multistart = lambda locally_optimize: memetic(
```

Black-box optimization

- heavy objective;
- non-differentiable:
 - complex computer simulations (e.g. aero-dynamics);
- so multi-modal, gradient does not have sense:
 - extremely deep networks (e.g. recurrent networks);

Surrogate optimization

A type of black-box optimization.

Given known samples $\mathcal{O}^t = \{(\theta_i, L_i)\}_{i=1}^t$:

- fit regression (surrogate) model to $\{(\theta_i, L_i)\}_{i=1}^t$;
- find the most promising θ_{t+1} with conventional optimization methods;
- evaluate objective in θ_{t+1} ;
- $\mathcal{O}^{t+1} = \mathcal{O}^t \cup \{(\theta_{t+1}, L_{t+1})\}$;
- repeat.

Bayesian optimization

The most well-known black-box optimization method.

Surrogate model:

- must estimate: $P(y \mid D, x)$;
- usually Gaussian processes:
 - easy to handle (normal distribution everywhere);
 - computationally expensive $O(n^3)$, n - number of points.
- possible to use Random Forest or Boosting;
- the most promising point is defined by aquisition function:

$$a(x) \rightarrow \max$$

Aquisition functions

- f' - current minimum; D - observed values.

Probability of improvement:

$$a_{pi}(x) = \mathbb{E} [\mathbb{I}[f(x) > f'] \mid D] = \underbrace{\Phi(f', \mu(x), K(x, x))}_{GP}$$

Expected improvement:

$$a_{ei}(x) = \mathbb{E} [\max(0, f' - f(x)) \mid D] = \underbrace{(f' - \mu(x))\Phi(f', \mu(x), K(x, x)) + K(x, x)\phi(f', \mu(x), K(x, x))}_{GP}$$

Lower confidence bound:

$$a_{lcb}(x) = \mu(x) + \beta\sigma(x)$$

Variational optimization

$$\min_x f(x) \leq \mathbb{E}_{x \sim \pi_\psi} f(x) = J(\psi)$$

- π_ψ - search distribution;

$$\nabla J(\psi) = \nabla \mathbb{E}_{x \sim \pi_\psi} f(x) =$$

$$\int_x f(x) (\nabla \pi(x | \psi)) \frac{\pi(x | \psi)}{\pi(x | \psi)} dx =$$

$$\int_x f(x) \nabla \log \pi(x | \psi) \pi(x | \psi) dx =$$

$$\mathbb{E}_{x \sim \pi_\psi} f(x) \nabla \log \pi(x | \psi) \approx$$

$$\frac{1}{n} \sum_{i=1}^n f(x_i) \nabla \log \pi(x_i | \psi)$$

Evolution Strategies

Evolution Strategies - is a subset of variational optimization (?).

- Repeat:
 - sample $\{x_i\}_{i=1}^n$ from $\pi(\cdot \mid \psi)$;
 - evaluate $f_i = f(x_i)$;
 - compute:

$$\nabla J(\psi) \approx \frac{1}{n} \sum_{i=1}^n f_i \nabla \log \pi(x_i \mid \psi)$$

- update ψ , e.g.

$$\psi \leftarrow \text{adamax}(\nabla J(\psi))$$

Evolution Strategies

- d - dimensionality of the problem;
- π_ψ - Gaussian:
 - $\dim(\psi) = O(d^2)$: covariance matrix and mean vector;
 - $O(d^3)$ operations per step;
- π_ψ - Gaussian with independent components:
 - $|\psi| = O(d)$: diagonal covariance matrix and mean vector;
 - $O(d)$ operations per step;
- π_ψ - scaled normal:
 - $|\psi| = O(d)$: variance σ and mean vector;
 - $O(d)$ operations per step;
 - less samples for estimating gradient.

Summary

Summary

Known your optimization algorithms:

- differentiable \Rightarrow gradient methods;
- super heavy objective \Rightarrow Bayesian;
- non-differentiable \Rightarrow Variational Optimization;
- wierd model \Rightarrow evolutionary optimization.

Don't blindly follow this.

- Bottou, L., 2012. Stochastic gradient descent tricks. In Neural networks: Tricks of the trade (pp. 421-436). Springer Berlin Heidelberg.
- Kingma, D. and Ba, J., 2014. Adam: A method for stochastic optimization. arXiv preprint arXiv:1412.6980.
- Zeiler, M.D., 2012. ADADELTA: an adaptive learning rate method. arXiv preprint arXiv:1212.5701.

- Fletcher, R., 2013. Practical methods of optimization. John Wiley Sons.

References, gradient-free

- Back, T., 1996. Evolutionary algorithms in theory and practice: evolution strategies, evolutionary programming, genetic algorithms. Oxford university press.
- Kennedy, J., 2011. Particle swarm optimization. In Encyclopedia of machine learning (pp. 760-766). Springer US.
- Wierstra, D., Schaul, T., Glasmachers, T., Sun, Y., Peters, J. and Schmidhuber, J., 2014. Natural evolution strategies. Journal of Machine Learning Research, 15(1), pp.949-980.
- Snoek, J., Larochelle, H. and Adams, R.P., 2012. Practical bayesian optimization of machine learning algorithms. In Advances in neural information processing systems (pp. 2951-2959).